

What is Julia ?

- High level, high performance dynamic programming language for technical computing
- Facilitates a distributed parallel execution
- Multiple dispatch, dynamic type system, macros and metaprogramming facilities are few of its main features
- Has been benchmarked for serial algorithms against R and Python

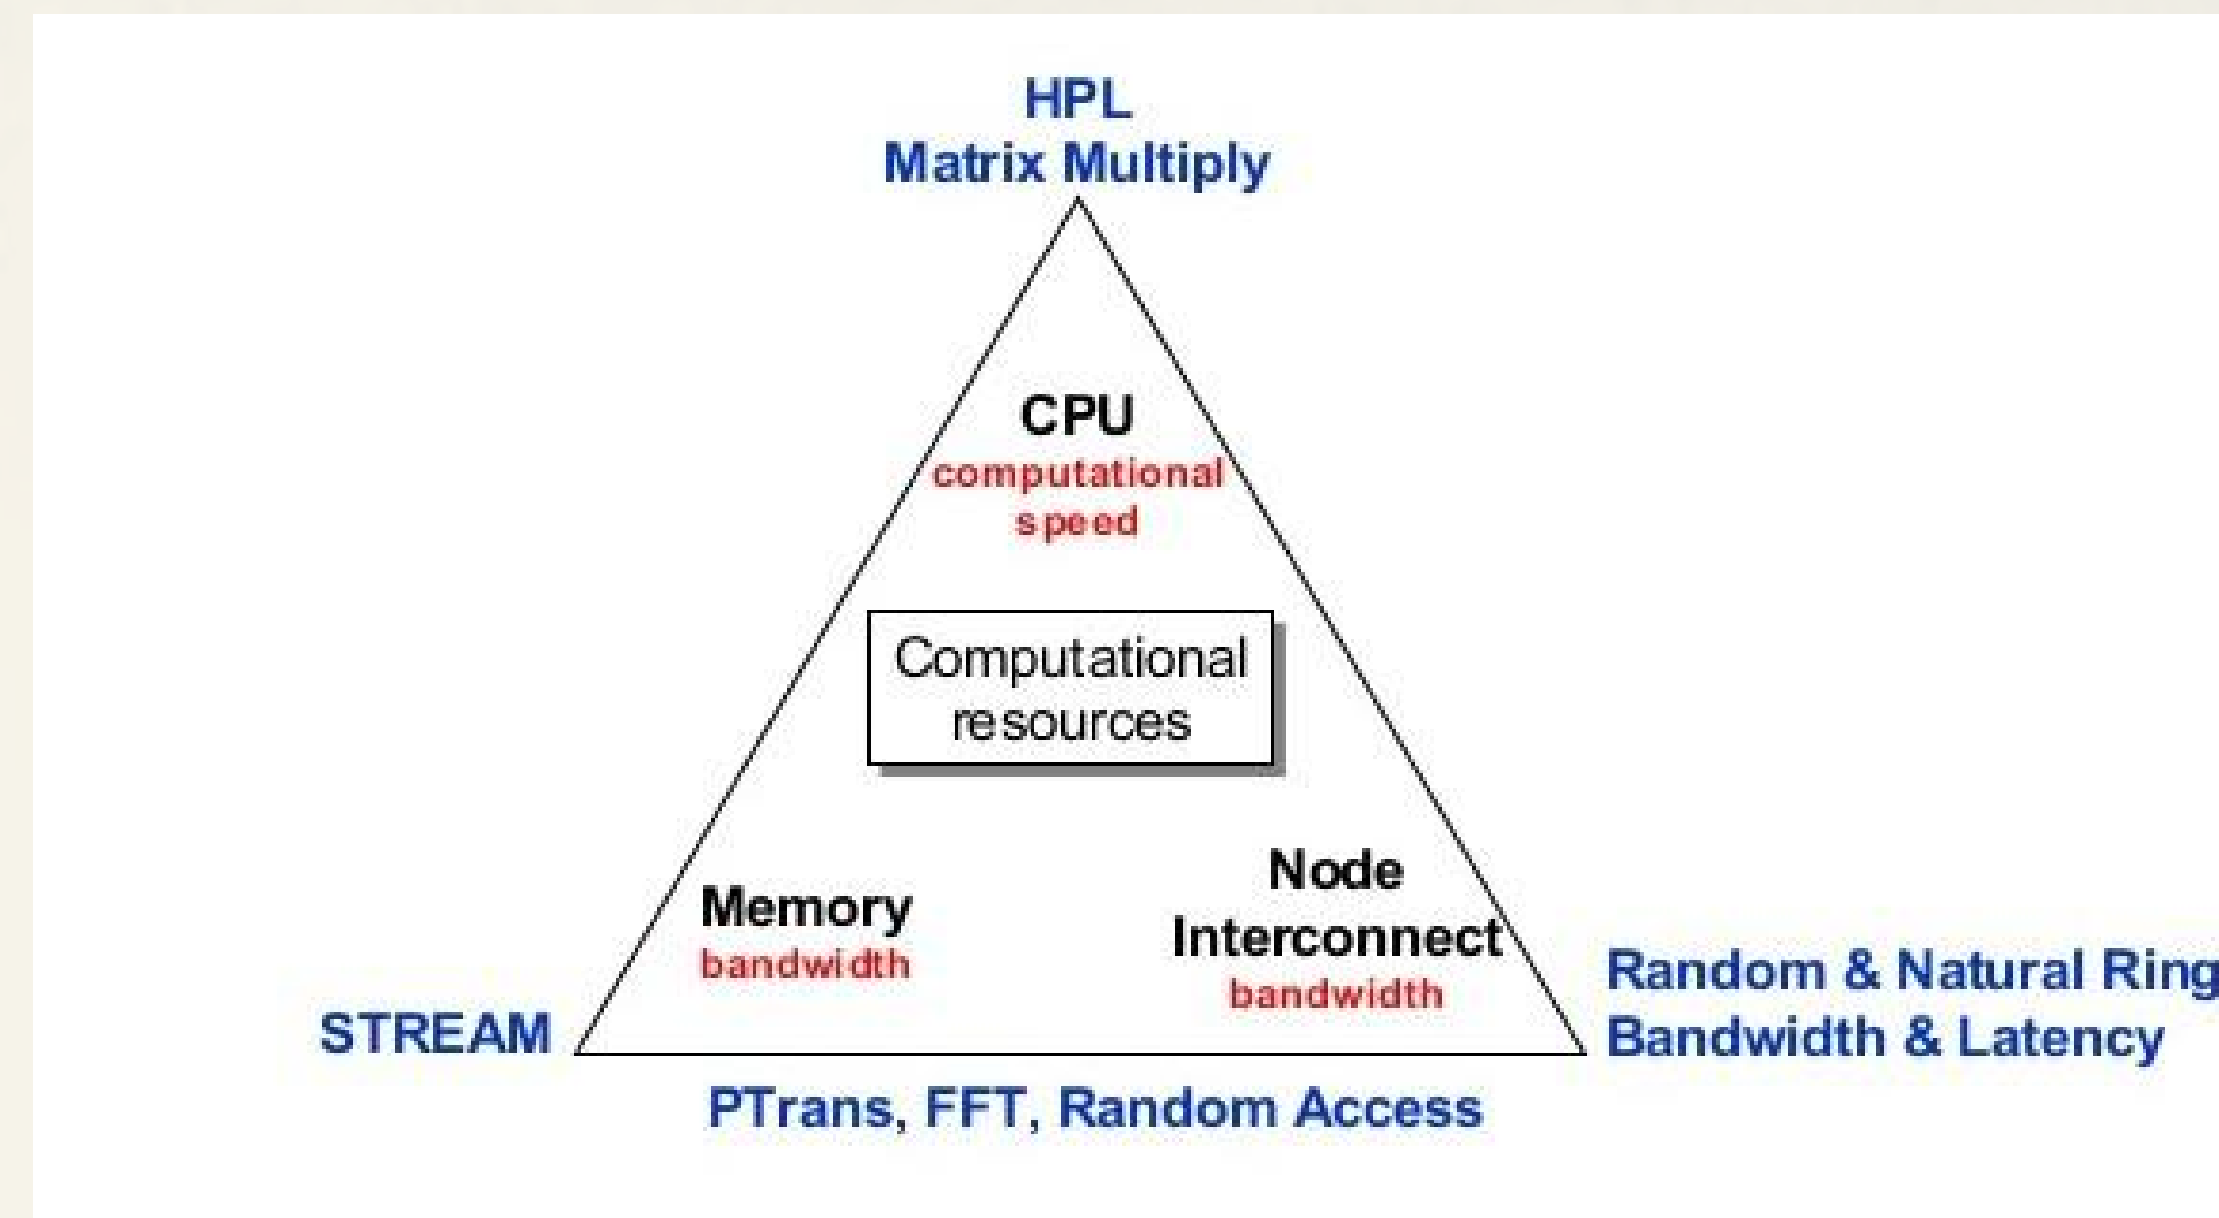
How does it work ?

- Multiprocessor environment - One master process and a team of workers
- Unlike MPI, communication is “one-sided”, ie., no “msg_send” or “msg_rcv”
- **Remote Calls** - Request by one process to call a certain function with some arguments on another process
- **Remote Ref** - Refers to an object stored on any process
- **fetch** - Explicit data transfer
- **@spawnat** - Evaluate an expression at specified processor
- Example -
julia> r = remotecall(2, rand, 2, 2)
julia> fetch(r)
julia> s = @spawnat 2 1 .+ fetch(r)
- **DArray** - Distributed array. ‘localpart’ and ‘localindexes’ accessible on a process
- **SharedArray** - Each process has access to entire array

Limitations

- Calls like MPI_Bcast, MPI_Reduce, MPI_SendRecv don't exist in Julia as all communication is one-sided
- Julia implementation does not take advantage of high-speed low-latency communication hardware such as InfiniBand interconnects in a cluster
- Julia opens a TCP port between every pair of processes that exchange data unlike MPI
- Bug: DArray memory is not fully garbage collected. This issue has been reported to Julia

HPCC Benchmarks

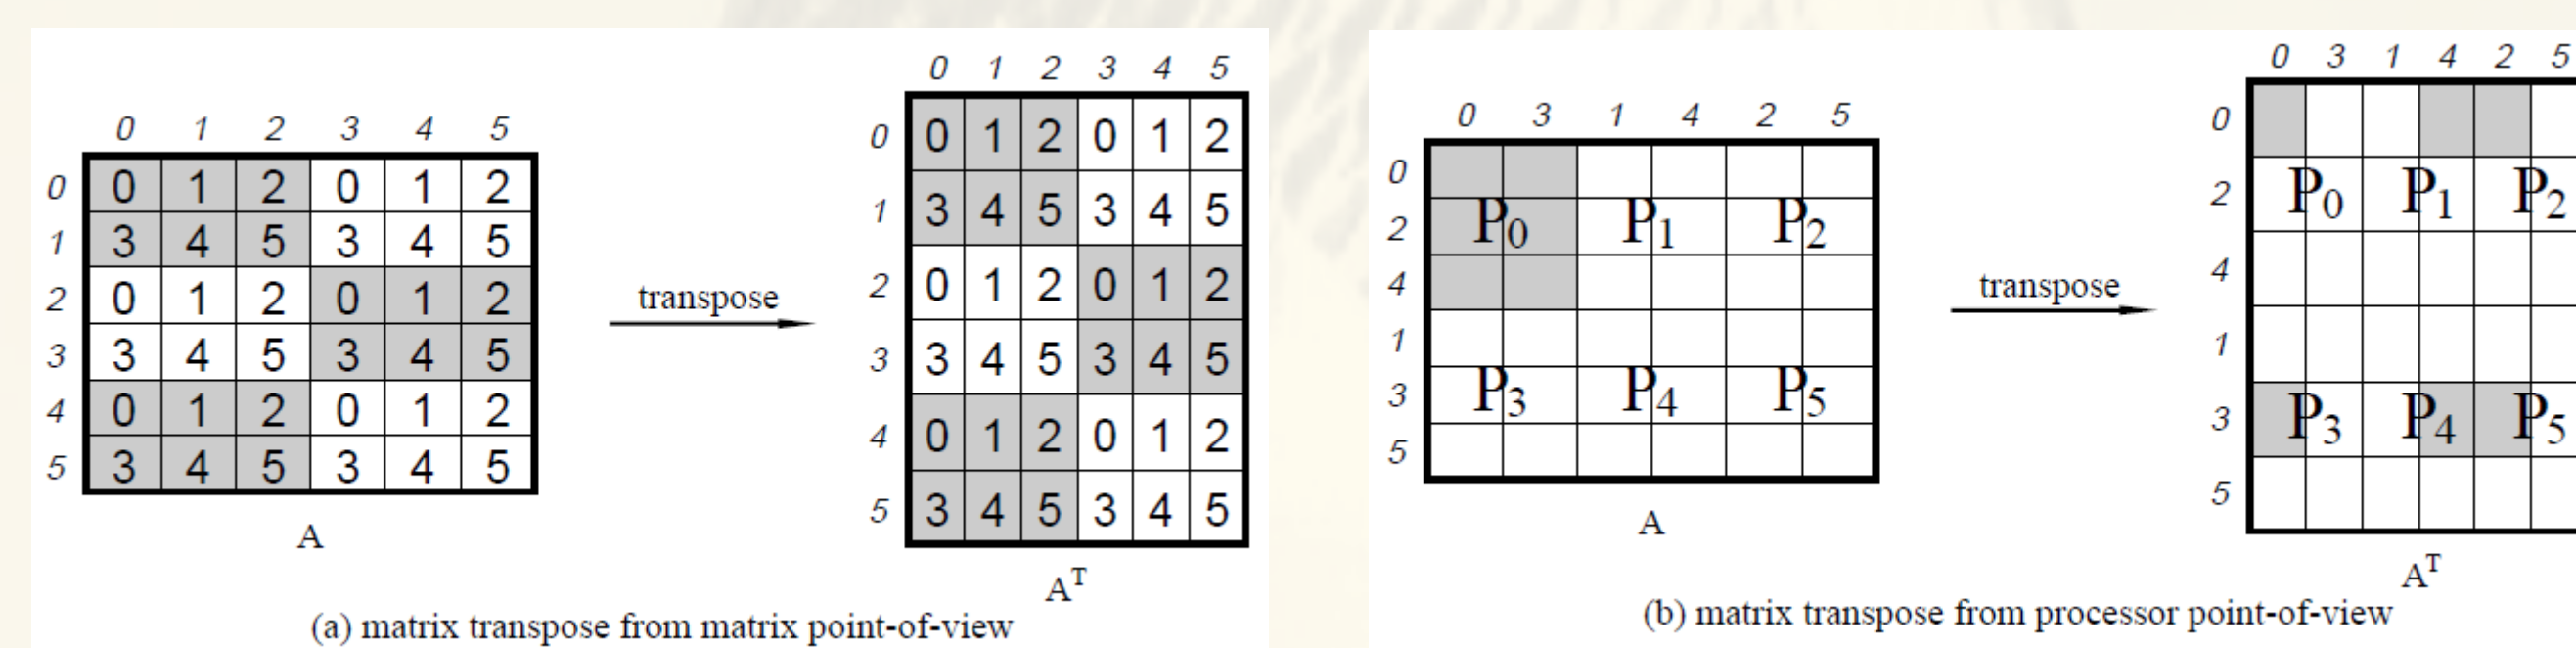


STREAM

- Measures sustainable memory bandwidth in GB/s.
- Measures the computation rate for vector kernels.
- Size of dataset should be greater than the sum of all last level caches available.
- Jinx L3 cache size is 8192 KB. We run STREAM on 16 cores, so we have used arrays containing 64 million elements each.
- Kernels-
 - “Copy” measures transfer rates in the absence of arithmetic.
 - “Scale” adds a simple arithmetic operation.
 - “Sum” adds a third operand to allow multiple load/store ports on vector machines to be tested.
 - “Triad” allows chained/overlapped/fused multiply/add operations.

PTRANS

- Parallel matrix transpose
- It is a test of the total communications capacity of the network.
- Measures rate of data transfer from multiprocessor memory and exercises communications where pairs of processors exchange large messages simultaneously.
- Ideas:
 - Matrix distributed over P X Q processor template
 - Block scattered data distribution
 - Non-blocking, point-to-point communication




Random Access

- Performance metric - GUPS
- It denotes number of memory locations that can be randomly updated in one second
- It profiles both memory bandwidth and node-interconnect bandwidth
- Implementation -
 - DArray across processors of total size 2^n
 - Each processor generates a stream of random 64-bit ints of total size 2^{n+2}
 - From each 64-bit int, highest n bits determine memory location
 - Number at the location xor-ed with 64-bit int
- Two additional variables -
 - Acceptable error
 - Bucket of updates

Parallel FFT




- Performance metric - MFLOPS
- It profiles both memory bandwidth and node-interconnect bandwidth
- Implementation -
 - Bit-reversed sequence of input array distributed across processes (DArray)
 - Each processor performs FFT on local chunk
 - log P rounds of inter-node communication as per Cooley-Tukey algorithm



Results and Analysis

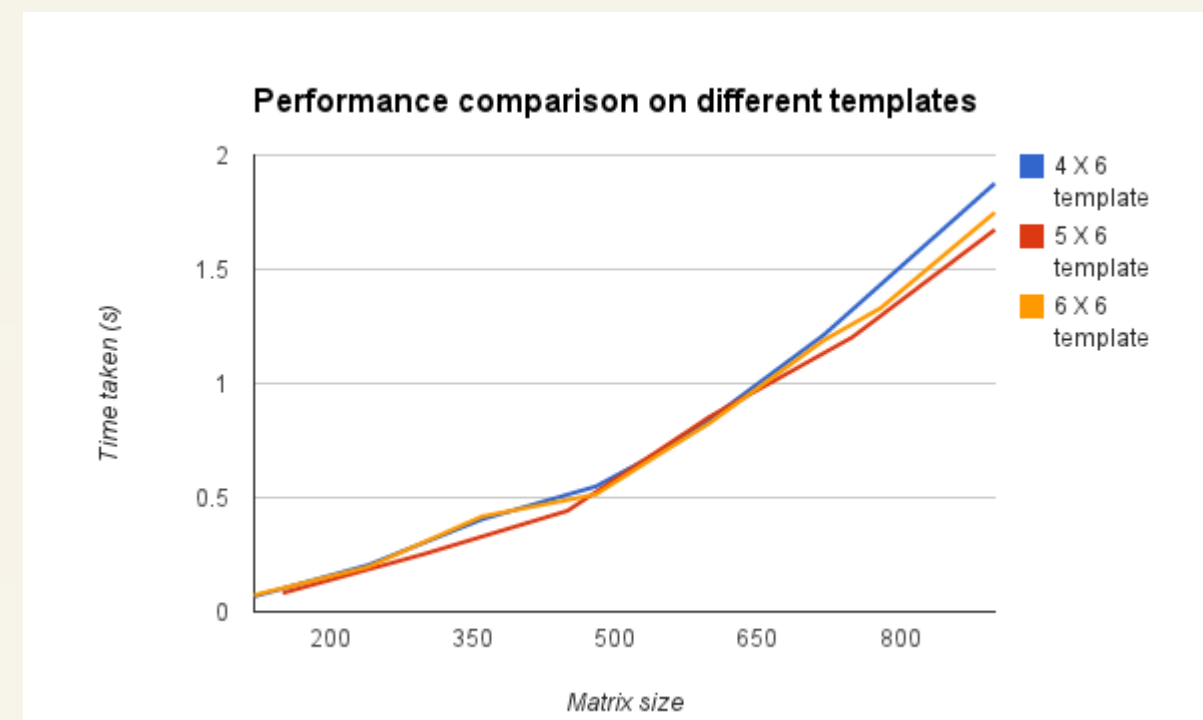
STREAM

- Performance of Julia comparable with C/OpenMP
- Sharing of L3 cache in case of single node
- Separate L3 cache in multi-node, less contention
- Initial increase in performance due to efficient resource utilization
- Less work than the number of workers leads to more memory latency




PTRANS

- Time to transpose increases as matrix size increases
- Similar performance for different P
- Time to transpose decreases as Q decreases



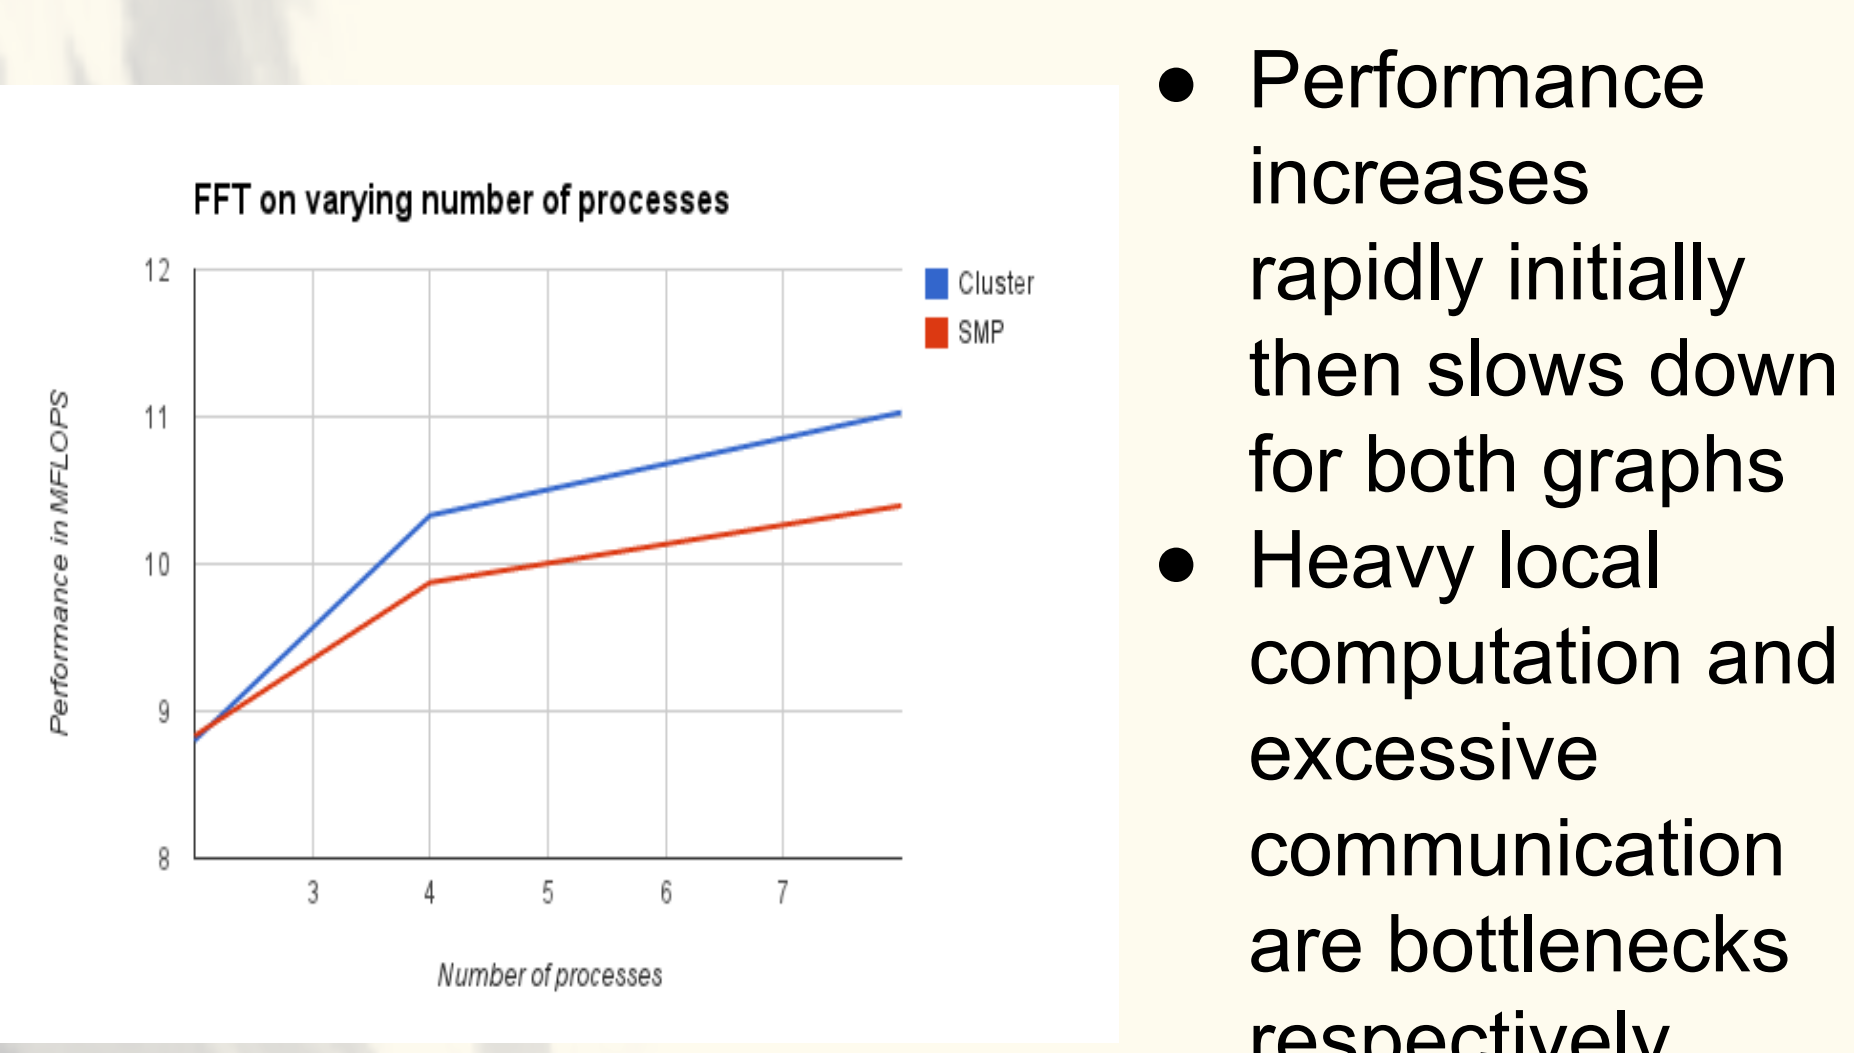

72 workers		64 workers		48 workers	
6 X 16	0.94 s	4 X 16	0.79 s	4 X 12	0.69 s
8 X 12	0.79 s	8 X 8	0.74 s	6 X 8	0.60 s
12 X 8	0.82 s			8 X 6	0.57 s
				12 X 4	0.59 s

Random Access



- As nodes increase, GUPS increase
- Even though 1% error is acceptable, Julia gives 0% failure rate everytime

Parallel FFT



- Performance increases rapidly initially then slows down for both graphs
- Heavy local computation and excessive communication are bottlenecks respectively